

Machine Learning – COMS3**007**



# Linear Regression

Benjamin Rosman

Based heavily on course notes by  
Chris Williams and Victor Lavrenko,  
Amos Storkey, Eric Eaton, and Clint  
van Alten

# Regression

- Data  $X = \{x^{(0)}, \dots, x^{(n)}\}$ , where  $x^{(i)} \in \mathbb{R}^d$
- Labels  $y = \{y^{(0)}, \dots, y^{(n)}\}$ , where  $y^{(i)} \in \mathbb{R}$
- Want to learn function  $y = f(x, \theta)$  to predict  $y$  for a new  $x$



# Linear regression – model

- Assume **model**:  $y = a + bx + \eta$ 
    - $\eta$  is Gaussian noise (don't model explicitly)
  - Higher dimensions:
    - $y = f(x, \theta) = \sum_{i=0}^d \theta_i x_i$
    - $y = \theta^T x$
  - Note: we call them weights  $w$  and parameters  $\theta$  interchangeably
  - This is linear regression: the model is **linear in the parameters**
- What is this noise? Where might it come from?
- Treat  $0^{th}$  dimension of  $x$  as 1 (i.e. the “intercept”)

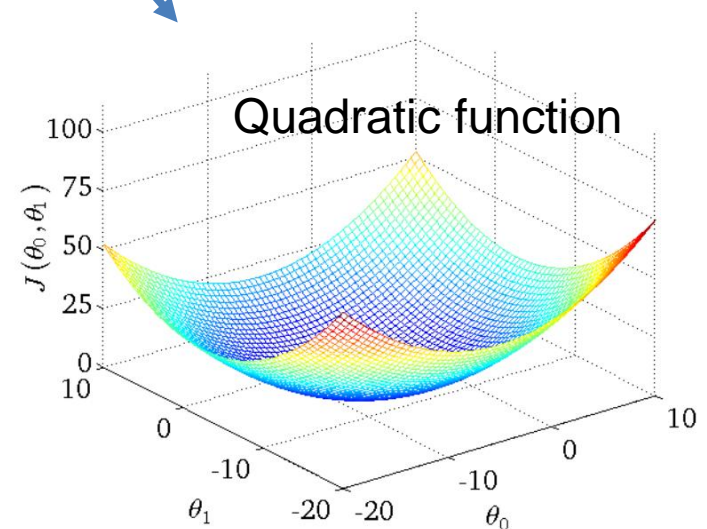

# Linear regression – cost function

- Infinitely many choices of  $\theta$ 
  - Learning = finding the best ones
- To choose among them, we need a cost function

$$E(\theta) = \frac{1}{2} \sum_{i=1}^n (y^{(i)} - f(x^{(i)}, \theta))^2$$


The  $\frac{1}{2}$  coefficient may instead be a 1, or a  $(1/2n)$ . It doesn't matter, as long as it is consistent.

It can also be labelled  $J()$  instead of  $E()$




# Linear regression – cost function

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



Error  
(function of the parameters  $\theta_0, \theta_1$ )




“weight space”


Every point here defines a regression line in the original problem

# Linear regression – cost function

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )




(function of the parameters  $\theta_0, \theta_1$ )




# Linear regression – cost function

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



(function of the parameters  $\theta_0, \theta_1$ )



# Basis functions



These basis functions  
are the features: good  
features are needed  
for regression to work!

- What makes it linear?
  - $y = f(x, \theta) = \sum_{i=0}^d \theta_i x_i$
  - Linear in  $\theta$ , NOT  $x$
- So: we can use different functions of  $x$ 
  - “Basis functions”  $\phi_j(x)$
  - Still linear regression
- Example: let  $x \in R^3$ , i.e.  $x = (x_1, x_2, x_3)^T$ 
  - Possible basis functions:
    - $x_1, x_1^5, x_1 x_2, x_3^2 x_2, \sin(x_2), \log(x_3), e^{-\frac{1}{2\sigma^2}(x_1 - \mu)^2}, \dots$
- Rewrite:  $y = f(x, \theta) = \sum_{i=0}^d \theta_i \phi_i(x)$



# Basis functions

- How to choose basis functions  $\phi_j(x)$ ?
  - Assumptions about the data
  - Try as many as possible
- Polynomial basis functions:
  - $\phi_j(x) = x^j$
  - Can include cross-terms
    - e.g.  $x_3^2 x_2$
  - **Global**: any change in  $x$  affects all basis functions
- Gaussian basis functions:
  - Radial basis functions (RBF)
  - $\phi_j(x) = e^{-\frac{1}{2\sigma^2}(x-\mu_j)^2}$
  - **Local**: change in  $x$  affects nearby basis functions
  - $\mu_j$  = location,  $\sigma$  = scale/width



# Design matrix

- Structure data in a design matrix  $\mathbf{X}$
- Let there be  $n$  data points (vectors):  $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(n)}$
- And  $d+1$  basis functions:  $\phi_0, \phi_1, \dots, \phi_d$

- Design matrix  $\Phi = \begin{bmatrix} \phi_0(\mathbf{x}^{(1)}) & \phi_1(\mathbf{x}^{(1)}) & \dots & \phi_d(\mathbf{x}^{(1)}) \\ \phi_0(\mathbf{x}^{(2)}) & \phi_1(\mathbf{x}^{(2)}) & \dots & \phi_d(\mathbf{x}^{(2)}) \\ \vdots & \vdots & \ddots & \vdots \\ \phi_0(\mathbf{x}^{(n)}) & \phi_1(\mathbf{x}^{(n)}) & \dots & \phi_d(\mathbf{x}^{(n)}) \end{bmatrix}$
- So far, we've had:  $\phi_0(\mathbf{x}) = 1, \phi_1(\mathbf{x}) = x_1, \phi_2(\mathbf{x}) = x_2, \text{ etc...}$
- Now we can work with the data in matrix form!

# Closed form solution

- How do we learn the function?
- We want to minimise the error:

- $E(\theta) = \frac{1}{2} \sum_{i=1}^n (y^{(i)} - f(x^{(i)}, \theta))^2$

- Rewrite:  $E(\theta) = \frac{1}{2} (\mathbf{y} - \mathbf{X}\theta)^2$

Vector notation  
with design matrix

- $E(\theta) = \frac{1}{2} (\mathbf{y} - \mathbf{X}\theta)^T (\mathbf{y} - \mathbf{X}\theta)$

- $E(\theta) = \frac{1}{2} (\mathbf{y}^T \mathbf{y} - 2\theta^T \mathbf{X}^T \mathbf{y} + \theta^T \mathbf{X}^T \mathbf{X} \theta)$

- For minimum: differentiate wrt  $\theta$  and set to zero


- $\theta = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$

The Moore-Penrose  
pseudo-inverse

# Example

- True function:  $f(x) = 0.2x^3 - 0.8x^2 - x + 1 + \eta$ 
  - Unknown to algorithm
- Training data:

$x$	$y$
0	1.1
1	-0.7
2	-2.6
3	-3.7
4	-2.9
5	1



# Example

- Choose model:  $f(x, \theta) = \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3$
- Error:  $E(\theta) = \frac{1}{2} \sum_{i=1}^n (y^{(i)} - f(x^{(i)}, \theta))^2$

Design matrix  $X =$


$x$	$y$
0	1.1
1	-0.7
2	-2.6
3	-3.7
4	-2.9
5	1

1	$x$	$x^2$	$x^3$
1	0	0	0
1	1	1	1
1	2	4	8
1	3	9	27
1	4	16	64
1	5	25	125

- $\theta = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} = [1.09, -1.3, -0.64, 0.18]^T$   
 $f(x) = 1 - x - 0.8x^2 + 0.2x^3$


# Example

- $\theta = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} = [1.09, -1.3, -0.64, 0.18]^T$



# Gradient descent

- Closed form solutions are not always appropriate
  - Can be slow: computing  $(\mathbf{X}^T \mathbf{X})^{-1}$  is roughly  $O(n^3)$
  - Cannot learn incrementally (redo computation for a new data point)
  - Different error functions?
- Instead use iterative procedure: gradient descent (GD)
- Basic idea:
  - Choose initial  $\theta$
  - Until a minimum:
    - Update  $\theta$  to reduce  $J(\theta)$



# Gradient descent


Cost function =  $E = J$

- Initialise  $\theta$
- Repeat until convergence:

- $\theta_j \leftarrow \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$

- Simultaneous update for  $j = 0, \dots, d$

$0 < \alpha \leq 1$  is the learning rate,  
usually set quite small



Take a **step of size  $\alpha$**   
in the “**downhill**”  
**direction** (negative  
gradient)



# Gradient descent

- Consider point  $i$
- $J(\theta) = E(\theta) = \frac{1}{2} (y^{(i)} - f(x^{(i)}, \theta))^2$
- For  $f(x^{(i)}, \theta) = \sum_{j=0}^d \theta_j x_j^{(i)}$ ,
  - $\frac{\partial}{\partial \theta_j} J(\theta) = (f(x^{(i)}, \theta) - y^{(i)}) x_j^{(i)}$


Stop when  $\|\theta_{new} - \theta_{old}\|_2 < \epsilon$

- So:
- Initialise  $\theta$
- Repeat until convergence:
  - For each datapoint  $i$ :
  - $\theta_j \leftarrow \theta_j - \alpha (f(x^{(i)}, \theta) - y^{(i)}) x_j^{(i)}$
  - Simultaneous update for  $j = 0, \dots, d$

# Gradient descent example


$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )




We will be taking downhill steps

# Gradient descent example


$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



$$J(\theta_0, \theta_1)$$


(function of the parameters  $\theta_0, \theta_1$ )



# Gradient descent example


$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



$$J(\theta_0, \theta_1)$$


(function of the parameters  $\theta_0, \theta_1$ )



# Gradient descent example


$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



$$J(\theta_0, \theta_1)$$


(function of the parameters  $\theta_0, \theta_1$ )



# Gradient descent example


$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



$$J(\theta_0, \theta_1)$$


(function of the parameters  $\theta_0, \theta_1$ )



# Gradient descent example


$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



$$J(\theta_0, \theta_1)$$


(function of the parameters  $\theta_0, \theta_1$ )



# Gradient descent example


$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )






# Gradient descent example


$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



$$J(\theta_0, \theta_1)$$


(function of the parameters  $\theta_0, \theta_1$ )



# Gradient descent example


$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



$$J(\theta_0, \theta_1)$$


(function of the parameters  $\theta_0, \theta_1$ )



# Example


- $\theta_0 \leftarrow \theta_0 - \alpha(f(x, \theta) - y^{(i)})$
- $\theta_1 \leftarrow \theta_1 - \alpha(f(x, \theta) - y^{(i)})x$
- $\theta_2 \leftarrow \theta_2 - \alpha(f(x, \theta) - y^{(i)})x^2$
- $\theta_3 \leftarrow \theta_3 - \alpha(f(x, \theta) - y^{(i)})x^3$

$$f(x, \theta) = \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3$$




# The effect of $\alpha$

$\alpha$  too small




$\alpha$  too large



- May overshoot the minimum
- May fail to converge
- May even diverge

# Overfitting and underfitting



- We don't just care about matching the training data, we want to generalise
- Too few (or the wrong) basis functions underfits
- But, too many overfits!
- Can we automate this?




# Regularisation

- $\theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3$
- $\approx \theta_0 + \theta_1 x$  if  $\theta_2, \theta_3 \approx 0$
- Idea: penalise large parameter values
  - Change the cost function
- $E(\theta) = \underbrace{\frac{1}{2} \sum_{i=1}^n (y^{(i)} - f(x^{(i)}, \theta))^2}_{\text{Fit the data}} + \underbrace{\lambda \sum_{j=1}^d \theta_j^2}_{\text{regularise}}$
- $\lambda \geq 0$  controls amount of regularisation
- Note: do not regularise  $\theta_0$

While minimising  $E$ , we want to **fit the data** AND **not have large parameters**:  
So, it will only grow parameters if the fit becomes much better!

# The effect of $\lambda$

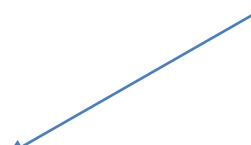
- $\lambda = 0$ : no regularisation
  - Linear regression as normal
- As  $\lambda$  increases (no upper bound):
  - It acts as an increasing force keeping parameters small unless really necessary
  - Tends to a straight line



# Closed-form **with regularisation**

- $\theta = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$

No regularisation  
on  $\theta_0$



- $\theta = \left( \mathbf{X}^T \mathbf{X} + \lambda \begin{bmatrix} 0 & 0 & 0 & \cdots & 0 \\ 0 & 1 & 0 & & 0 \\ 0 & 0 & 1 & & 0 \\ & \vdots & & \ddots & \vdots \\ 0 & 0 & 0 & \cdots & 1 \end{bmatrix} \right)^{-1} \mathbf{X}^T \mathbf{y}$


- Can derive this as before



# GD with regularisation

- Initialise  $\theta$
- Repeat until convergence:
  - $\theta_0 \leftarrow \theta_0 - \alpha(f(x^{(i)}, \theta) - y^{(i)})$
  - $\theta_j \leftarrow \theta_j - \alpha \left[ (f(x^{(i)}, \theta) - y^{(i)})x_j^{(i)} + \lambda\theta_j \right]$
- Simultaneous update for  $j = 0, \dots, d$

No regularisation  
on  $\theta_0$



# Recap

- Model
- Cost function
- Basis functions
- Design matrix
- Closed-form solution
- Gradient descent
- Regularisation