

Machine Learning – COMS3**007**

Neural Networks

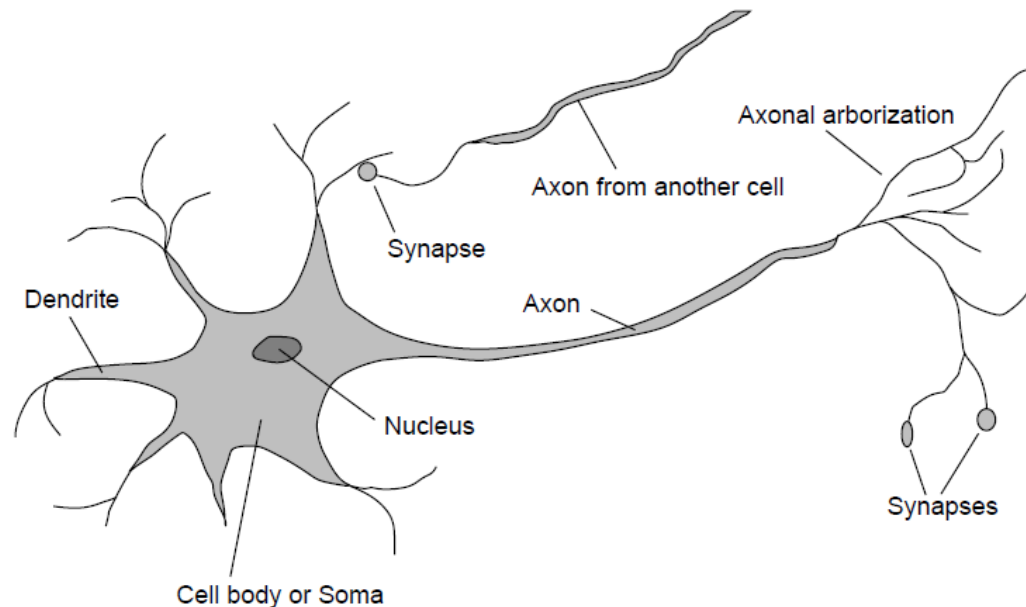
(Representations)

Benjamin Rosman

Based heavily on course notes by
Geoffrey Hinton, Chris Williams and
Victor Lavrenko, Amos Storkey, Eric
Eaton, and Clint van Alten

A biological perspective

- Interest in **how the brain works** (as a computing machine)
 - Human brain: 10^{11} neurons, 10^{14} synapses (connections)
 - Computers are much faster
 - But there are many more neurons
 - Neurons operate in parallel



A biological perspective

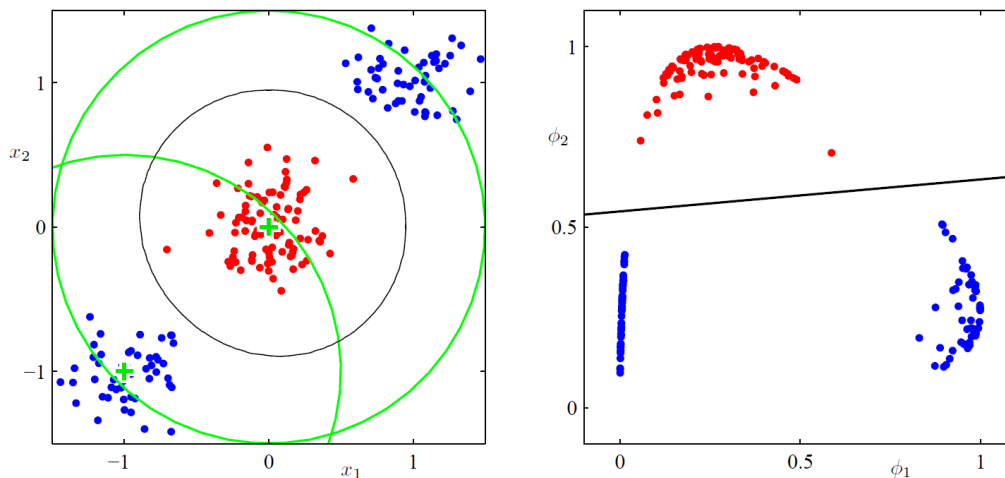
- Neurons send signals to each other
- Different parts of the brain do different things
- “One learning algorithm” hypothesis
 - Different areas learn in the same way
- Connection strength (synaptic weights) can change
 - Learning: perform useful computations
 - Strengths modulate the effects of signals between neurons
 - Excitatory vs inhibitory

Artificial neural networks

- Idealised models of how natural neurons work
- Allows us to study:
 - How the brain works
 - How to do computations in parallel rather than sequentially
 - How to solve complex problems

An engineering perspective

- We've spoken a lot about the importance of choosing the right features $\phi(x)$:

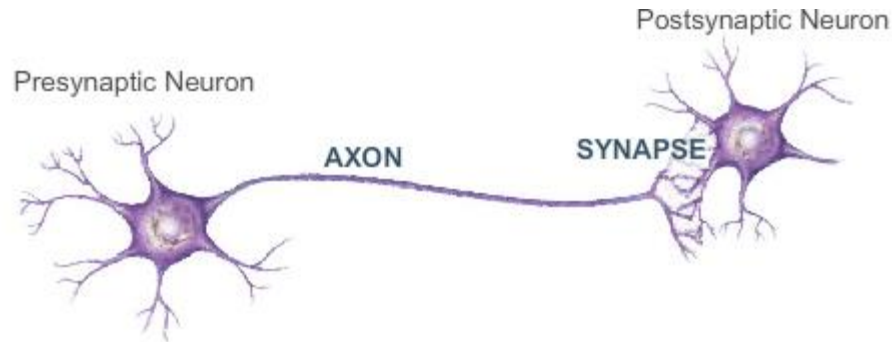


- Feature engineering (using domain expertise) has always been the **most important aspect**
 - But can't we learn these features as well as the weights?

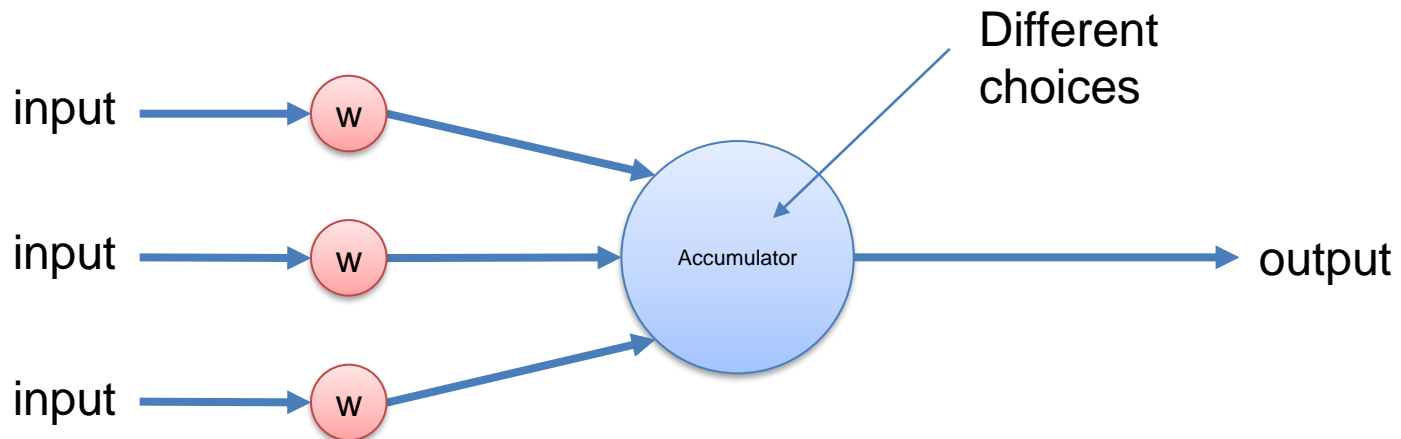
A brief history

- 1943: first artificial neuron: threshold logic unit
 - Warren McCulloch and Walter Pitts
- 1949: weights (synapses) adapt during learning
 - Donald Hebb
- 1957: the perceptron (Frank Rosenblatt)
- 1969: “Perceptrons” (Minsky and Papert)
- 1970s: the AI winter
- 1975: backpropagation – learning in networks (Paul Werbos)
- 1986: backprop rediscovered
 - David Rumelhart and Geoffrey Hinton
- 1980s – 2006: difficult to train networks
- 2006: new strategies for training networks
 - Geoffrey Hinton, Yann LeCun, Yoshua Bengio
- Since then: deep learning everywhere!

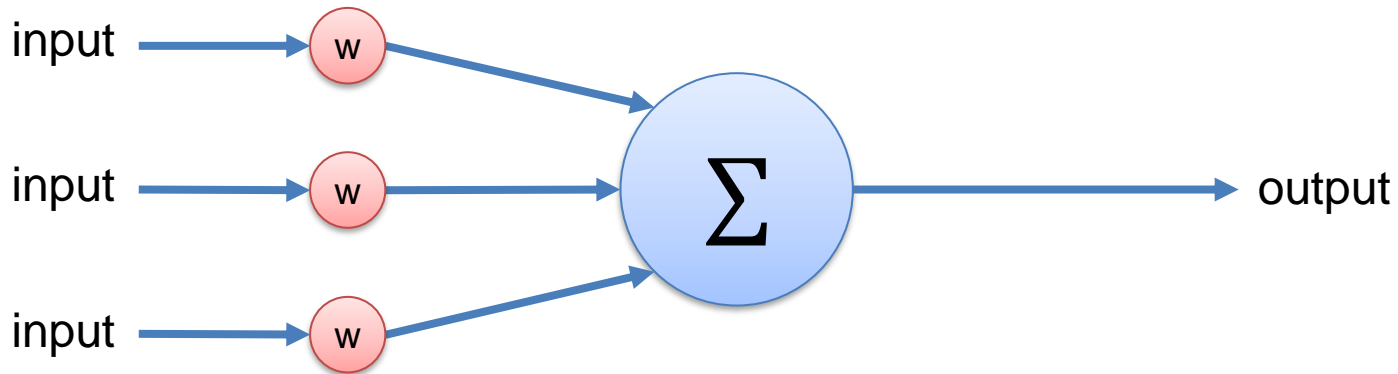
A simple neuron



- Artificial neuron:
 - Output y as a function of inputs x and weights w



Linear neurons




- Let one of these inputs be a bias $b = 1 = x_0$

- $y = \sum_{i=0}^d x_i w_i$
 - Sum over all input connections (including bias)
 - x_i : i^{th} input
 - w_i : Weight on i^{th} input

- But this is just **linear regression**!

Binary threshold neurons



- Use a step function (threshold function)

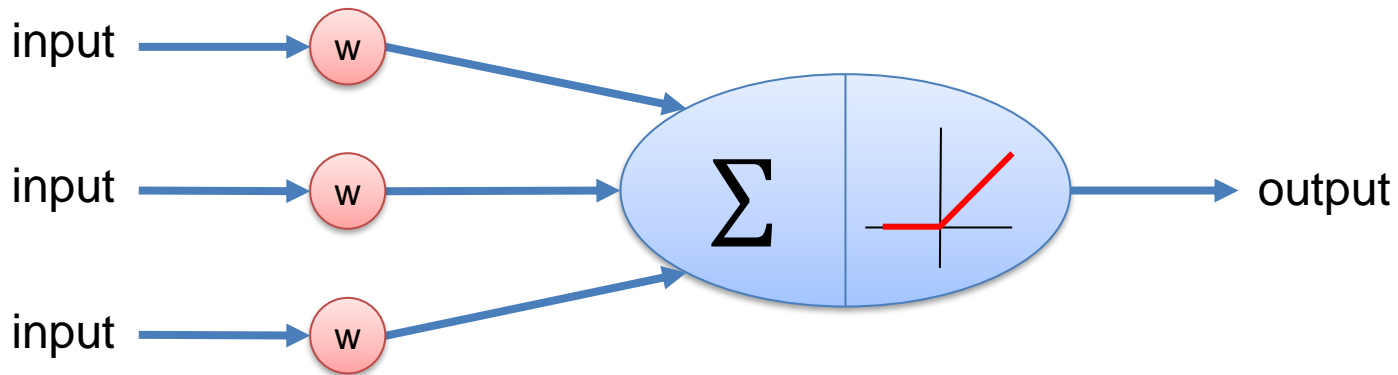
- $z = \sum_{i=0}^d x_i w_i$

- $y = \begin{cases} 1 & \text{if } z \geq \theta \\ 0 & \text{otherwise} \end{cases}$

Threshold (could be 0)

- But this is the **perceptron**!

Rectified linear neurons (ReLU)

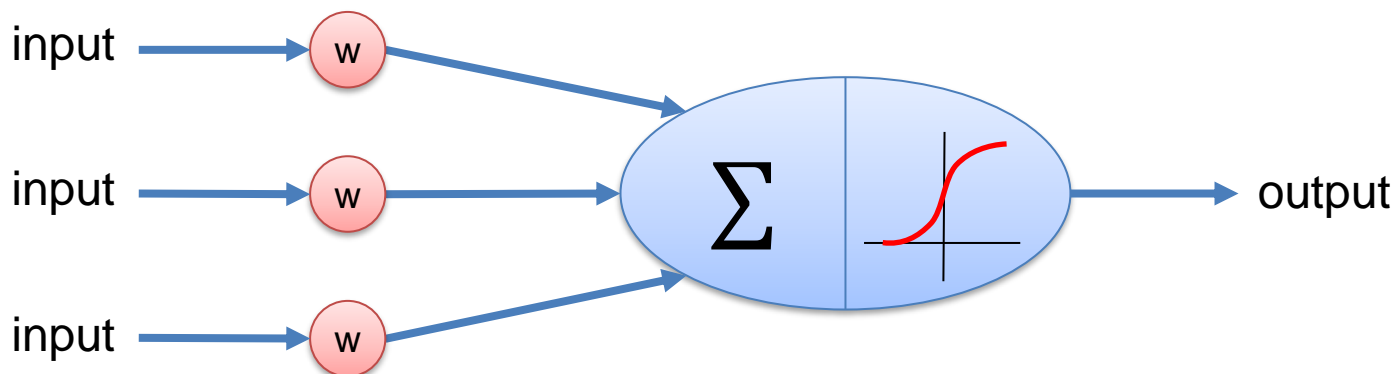


- Somewhere between a linear and threshold unit

- $z = \sum_{i=0}^d x_i w_i$

- $y = \begin{cases} \textcolor{red}{z} & \text{if } z \geq 0 \\ 0 & \text{otherwise} \end{cases}$

Sigmoid/logistic neurons



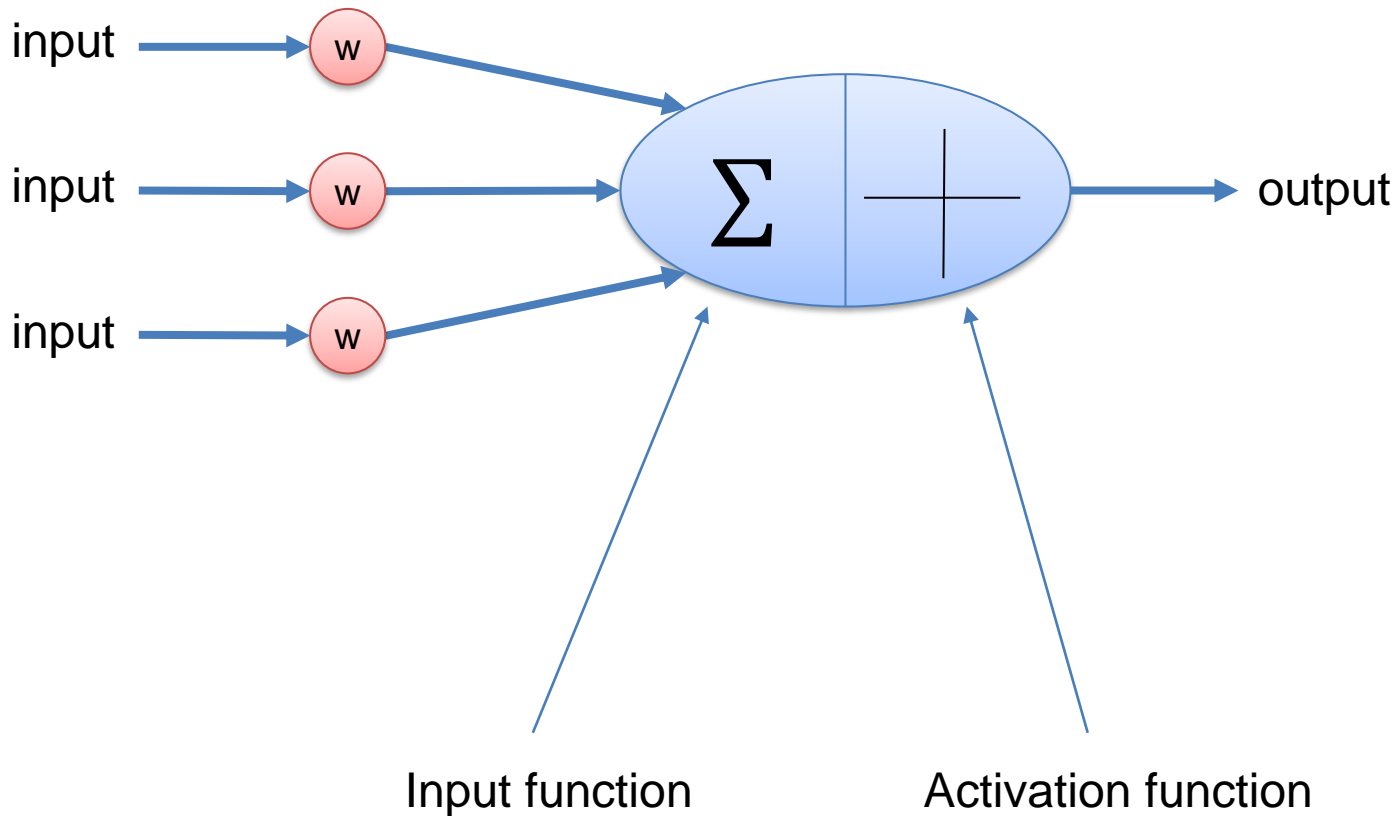
- Use a logistic function

- $z = \sum_{i=0}^d x_i w_i$

- $y = \frac{1}{1 + e^{-z}}$

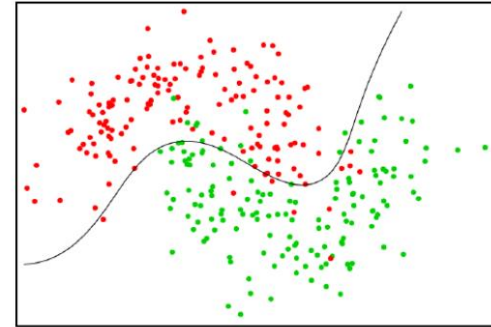
- But this is **logistic regression**!

Neuron anatomy



- Intuition: stack these neurons to learn features!

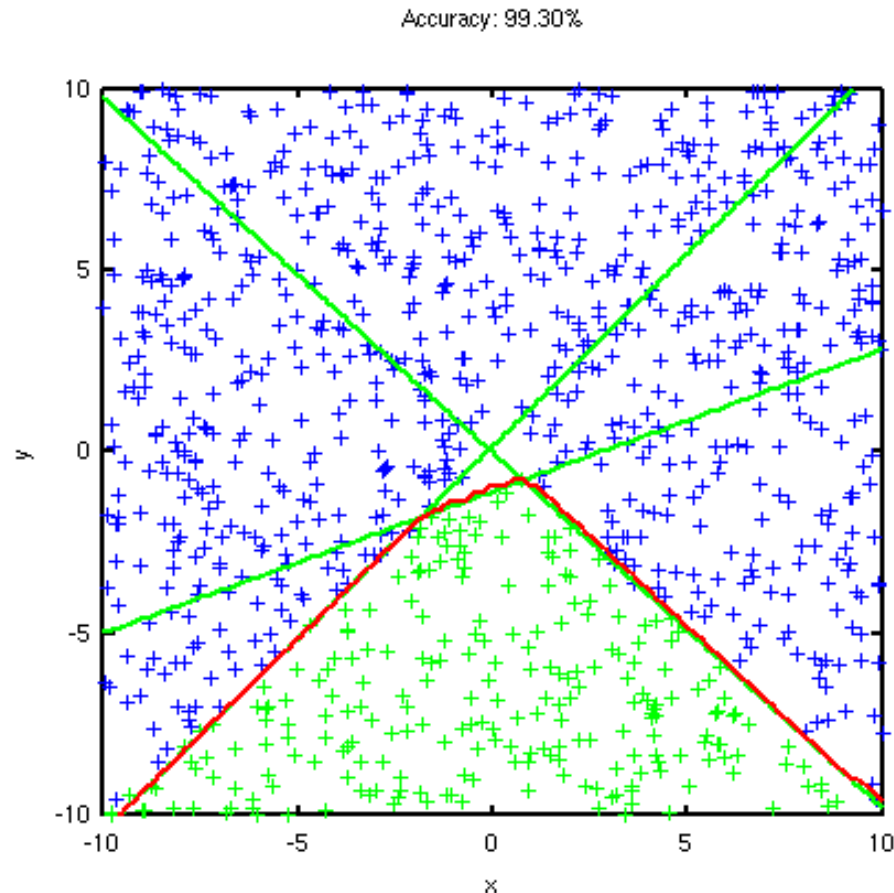
Learning features



- Logistic regression performs classification
 - Only with linear decision boundaries!
- To model more complicated systems, we need nonlinear class boundaries
- Use features $\phi(x)$ to get these nonlinear decision boundaries. But what features?
 - Ideally: features that apply to different parts of the input space
- Logistic regression model divides input space into two!
 - Let each feature $\phi_j(x)$ be a logistic regression model (neuron)

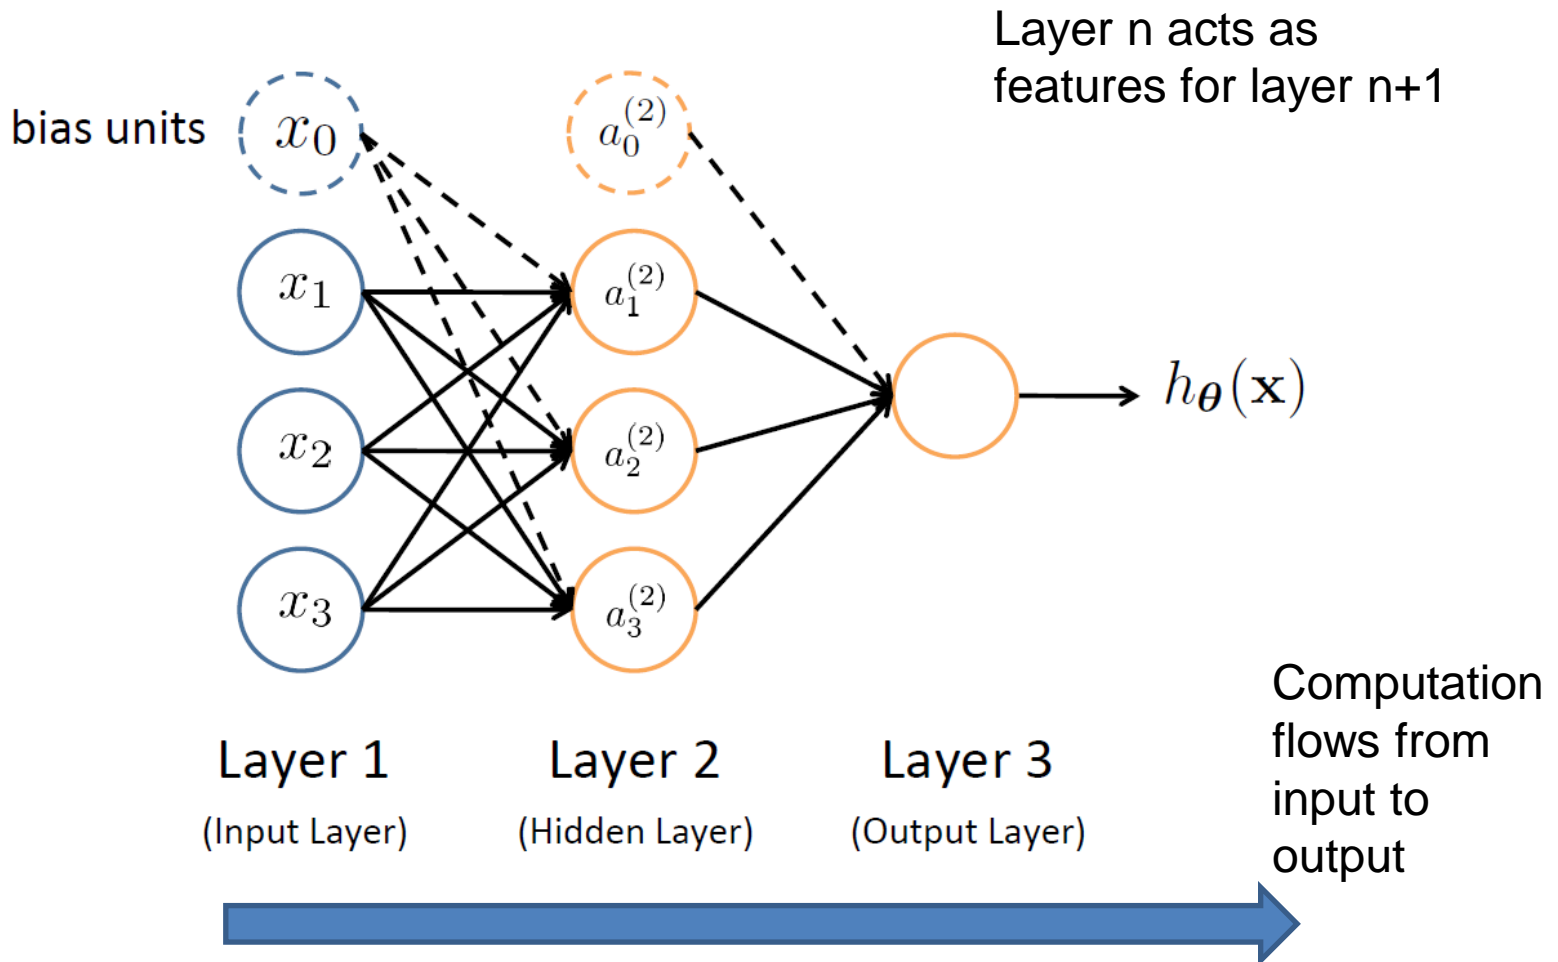
Learning features

- Logistic regression model divides input space into two!
 - Let each feature $\phi_j(x)$ be a logistic regression model (neuron)



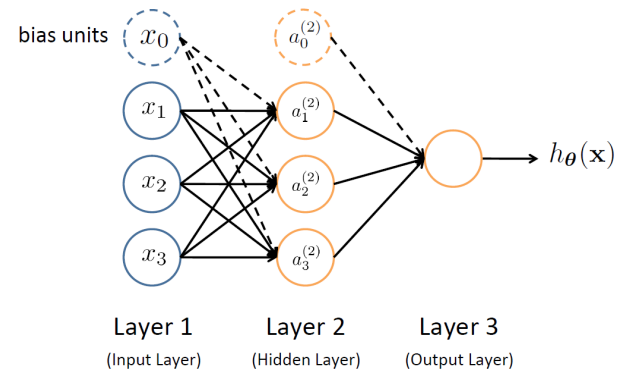
Stacking neurons

- Neurons arranged in layers



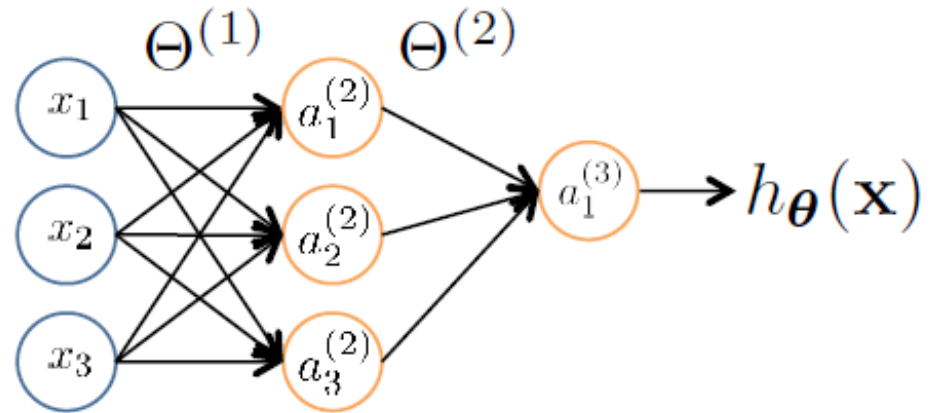
Feed-forward networks

Sometimes called multilayer perceptrons (MLPs)



- Input layers
 - Raw data
 - As provided by sensor measurements
- Feed-forward networks (most common)
 - **Outputs from one layer become inputs to the next**
- Working forward through the network:
 - Apply **input function** to compute total input
 - Usually just the sum of inputs
 - **Activation function** transforms input to final value
 - Usually nonlinear function
- Output layer: computation target

Computations



- $a_i^{(j)}$ = activation of unit i in layer j
- $\Theta^{(j)}$ = weight matrix: mapping from layer j to $j+1$

- $a_1^{(2)} = g \left(\Theta_{10}^{(1)} x_0 + \Theta_{11}^{(1)} x_1 + \Theta_{12}^{(1)} x_2 + \Theta_{13}^{(1)} x_3 \right)$
- $a_2^{(2)} = g \left(\Theta_{20}^{(1)} x_0 + \Theta_{21}^{(1)} x_1 + \Theta_{22}^{(1)} x_2 + \Theta_{23}^{(1)} x_3 \right)$
- $a_3^{(2)} = g \left(\Theta_{30}^{(1)} x_0 + \Theta_{31}^{(1)} x_1 + \Theta_{32}^{(1)} x_2 + \Theta_{33}^{(1)} x_3 \right)$

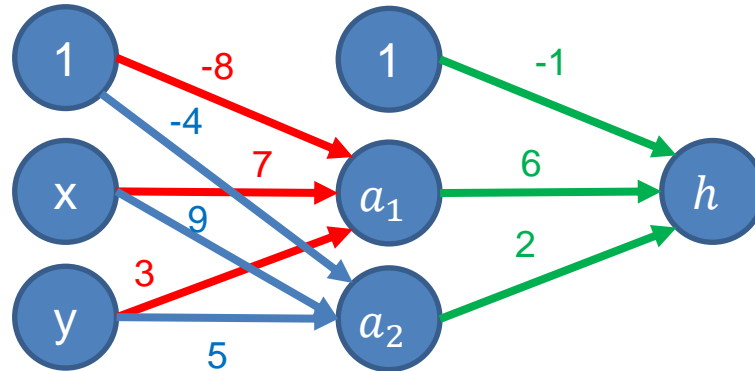
- $h_{\theta}(x) = a_1^{(3)} = g \left(\Theta_{10}^{(2)} a_0^{(2)} + \Theta_{11}^{(2)} a_1^{(2)} + \Theta_{12}^{(2)} a_2^{(2)} + \Theta_{13}^{(2)} a_3^{(2)} \right)$

- If s_j units in layer j , s_{j+1} in layer $j+1$: $\dim(\Theta^{(j)}) = s_{j+1} \times (s_j + 1)$

Activation
function


Bias = 1

Example




- Assume all activation functions are $g()$
- Output?
- $h = g(6a_1 + 2a_2 - 1)$
- $h = g(6g(7x + 3y - 8) + 2g(9x + 5y - 4) - 1)$

Example




- Weight matrix (layer 1) = $W^{(1)} = \begin{bmatrix} -8 & 7 & 3 \\ -4 & 9 & 5 \end{bmatrix}$
- Weight matrix (layer 2) = $W^{(2)} = \begin{bmatrix} -1 & 6 & 2 \end{bmatrix}$
- Assume all activation functions are ReLUs
- E.g. Input: $x = (1.5, -1)^T$
- Output?
 - Compute a_1, a_2 and then propagate forward to h

Example



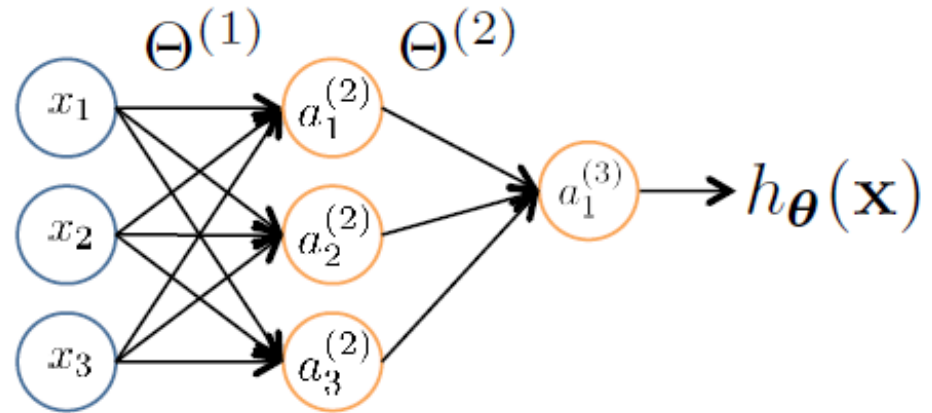
- $W^{(1)} = \begin{bmatrix} -8 & 7 & 3 \\ -4 & 9 & 5 \end{bmatrix}$, $W^{(2)} = \begin{bmatrix} -1 & 6 & 2 \end{bmatrix}$
- All activation functions are ReLUs
- Input: $x = (1.5, -1)^T \rightarrow$ augment to $x = (\textcolor{teal}{1}, 1.5, -1)^T$
- Total input to a_1 :
 - $\sum_{i=1}^3 w_{1i}^{(1)} x_i = (-8)(1) + (7)(1.5) + (3)(-1) = -0.5$
- Activation at a_1 (ReLU):
 - $-0.5 < 0$ so output $a_1 = 0$
- Total input to a_2 :
 - $\sum_{i=1}^3 w_{2i}^{(1)} x_i = (-4)(1) + (9)(1.5) + (5)(-1) = 4.5$
- Activation at a_2 (ReLU):
 - $4.5 > 0$ so output $a_2 = 4.5$

Example



- $W^{(1)} = \begin{bmatrix} -8 & 7 & 3 \\ -4 & 9 & 5 \end{bmatrix}$, $W^{(2)} = \begin{bmatrix} -1 & 6 & 2 \end{bmatrix}$
- All activation functions are ReLUs
- $a_1 = 0$
- $a_2 = 4.5$
- Augment with $a_0 = 1$: $a = (1, 0, 4.5)^T$
- Total input to h :
 - $\sum_{i=1}^3 w_i^{(2)} a_i = (-1)(1) + (6)(0) + (2)(4.5) = 8$
- Activation at h (ReLU):
 - $8 > 0$ so output $h = 8$

Vectorization




- $a_1^{(2)} = g \left(\Theta_{10}^{(1)} x_0 + \Theta_{11}^{(1)} x_1 + \Theta_{12}^{(1)} x_2 + \Theta_{13}^{(1)} x_3 \right) = g(z_1^{(2)})$
- $a_2^{(2)} = g \left(\Theta_{20}^{(1)} x_0 + \Theta_{21}^{(1)} x_1 + \Theta_{22}^{(1)} x_2 + \Theta_{23}^{(1)} x_3 \right) = g(z_2^{(2)})$
- $a_3^{(2)} = g \left(\Theta_{30}^{(1)} x_0 + \Theta_{31}^{(1)} x_1 + \Theta_{32}^{(1)} x_2 + \Theta_{33}^{(1)} x_3 \right) = g(z_3^{(2)})$
- $h_{\theta}(x) = g \left(\Theta_{10}^{(2)} a_0^{(2)} + \Theta_{11}^{(2)} a_1^{(2)} + \Theta_{12}^{(2)} a_2^{(2)} + \Theta_{13}^{(2)} a_3^{(2)} \right) = g(z_1^{(3)})$

Vectorized steps:

- $\mathbf{z}^{(2)} = \Theta^{(1)} \mathbf{x}$
- $\mathbf{a}^{(2)} = g(\mathbf{z}^{(2)})$
- Add $a_0^{(2)} = 1$
- $\mathbf{z}^{(3)} = \Theta^{(2)} \mathbf{a}^{(2)}$
- $h_{\theta}(\mathbf{x}) = \mathbf{a}^{(3)} = g(\mathbf{z}^{(3)})$





Example



- $W^{(1)} = \begin{bmatrix} -8 & 7 & 3 \\ -4 & 9 & 5 \end{bmatrix}$, $W^{(2)} = \begin{bmatrix} -1 & 6 & 2 \end{bmatrix}$
- All activation functions are ReLUs
- Input: $x = (1.5, -1)^T \rightarrow$ augment to $x = (\mathbf{1}, 1.5, -1)^T$
- Total input $z^{(2)}$ at layer 2:
 - $z^{(2)} = W^{(1)}x = \begin{bmatrix} -8 & 7 & 3 \\ -4 & 9 & 5 \end{bmatrix} \begin{bmatrix} 1 \\ 1.5 \\ -1 \end{bmatrix} = \begin{bmatrix} -0.5 \\ 4.5 \end{bmatrix}$
- Apply ReLU operator: $a^{(2)} = g(z^{(2)}) = (0, 4.5)^T$
- Augment with $a_0 = 1$: $a = (\mathbf{1}, 0, 4.5)^T$
- Total input $z^{(3)}$ at layer 3:
 - $z^{(3)} = W^{(2)}a = \begin{bmatrix} -1 & 6 & 2 \end{bmatrix} (1, 0, 4.5)^T = 8$
- Apply ReLU operator: $h = g(z^{(3)}) = 8$


Hidden neurons

- If each neuron produces a sigmoid:
- Two can combine to give a ridge:
(in the next layer)
- These can create more complex structures:



Architectures

- Can build many different architectures
- E.g.



- Let L be the number of layers, s_i = nodes on layer i
 - Hyperparameters
- Usually: $s_0 = d$ (# input features), $s_{L-1} = K$ (# classes)

Multiclass classification



Pedestrian



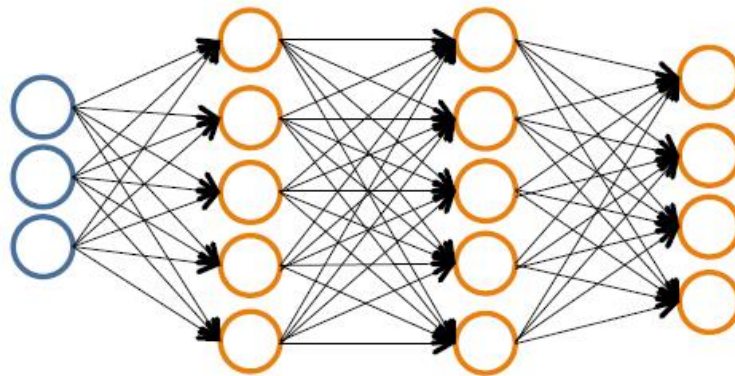
Car



Motorcycle



Truck



$$h_{\Theta}(\mathbf{x}) \in \mathbb{R}^K$$

We want:

$$h_{\Theta}(\mathbf{x}) \approx \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

when pedestrian

$$h_{\Theta}(\mathbf{x}) \approx \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$$

when car


$$h_{\Theta}(\mathbf{x}) \approx \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$$

when motorcycle

$$h_{\Theta}(\mathbf{x}) \approx \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$$

when truck

Multiclass classification



$$h_{\Theta}(\mathbf{x}) \in \mathbb{R}^K$$

We want:

$$h_{\Theta}(\mathbf{x}) \approx \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

when pedestrian

$$h_{\Theta}(\mathbf{x}) \approx \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$$

when car

$$h_{\Theta}(\mathbf{x}) \approx \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$$

when motorcycle

$$h_{\Theta}(\mathbf{x}) \approx \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$$

when truck

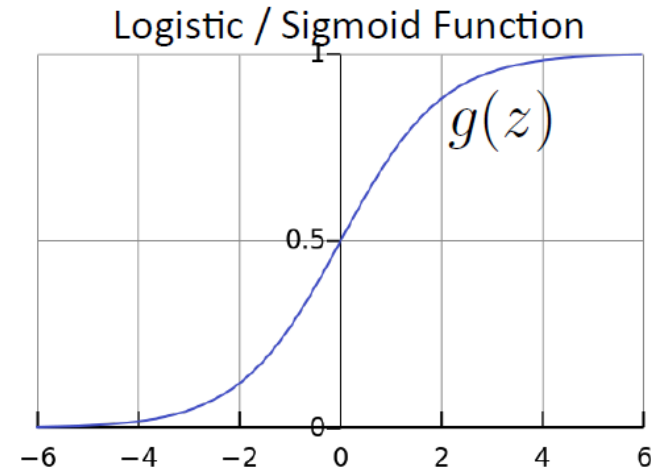
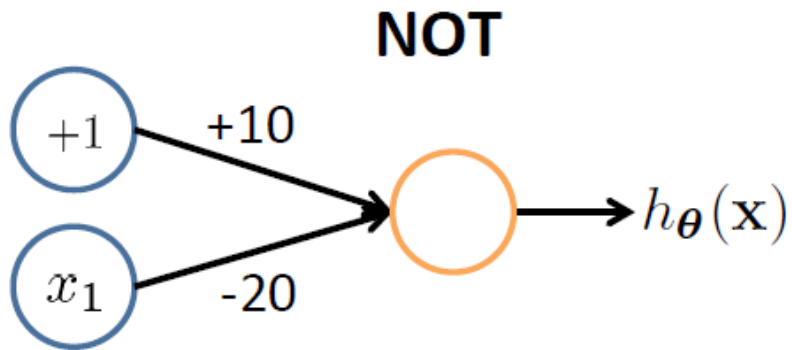
- Given data $\{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n)\}$
- Convert labels to **1-of-K representation**
 - E.g. $y_i = \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$ when motorcycle, $y_i = \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$ when car, etc.

Representation power

- Every Boolean function can be represented by a network with a single hidden layer
- Every bounded continuous function can be approximated with arbitrarily small error by a network with one hidden layer
- Any function can be approximated to arbitrary accuracy by a network with two hidden layers
- Neural networks are **universal function approximators**
- But:
 - A function being representable does not tell us how many hidden units would be required
 - May be exponential!
 - Nor how easily this can be learned!

Boolean functions – NOT

- Consider NOT
- $x_1 \in \{0,1\}$
- $y = \text{NOT } x_1$



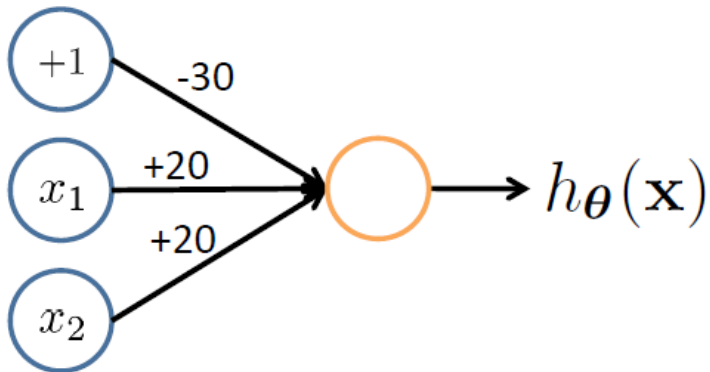
x_1	$h_{\theta}(x)$
0	$g(10) \approx 1$
1	$g(-10) \approx 0$

Boolean functions – AND



Simple example: AND

$$x_1, x_2 \in \{0, 1\}$$

$$y = x_1 \text{ AND } x_2$$




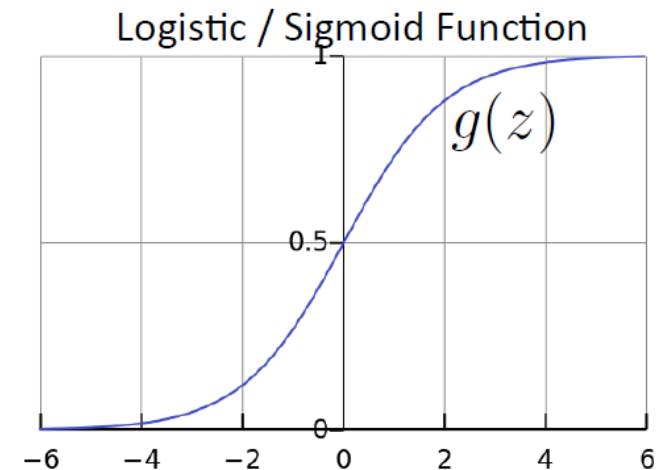
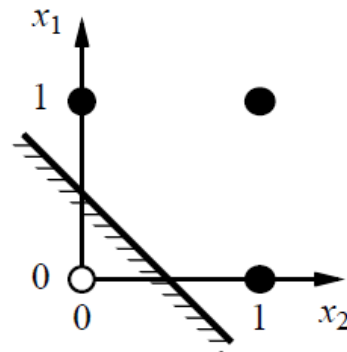
$$h_{\theta}(\mathbf{x}) = g(-30 + 20x_1 + 20x_2)$$



x_1	x_2	$h_{\theta}(\mathbf{x})$
0	0	$g(-30) \approx 0$
0	1	$g(-10) \approx 0$
1	0	$g(-10) \approx 0$
1	1	$g(10) \approx 1$

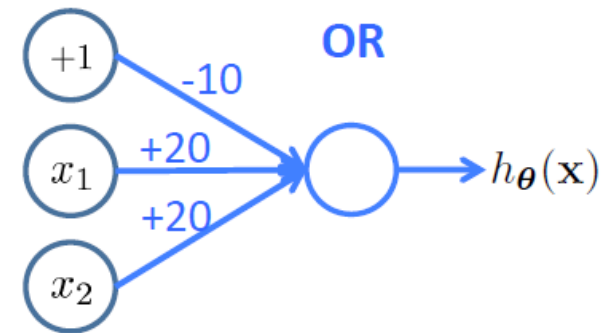

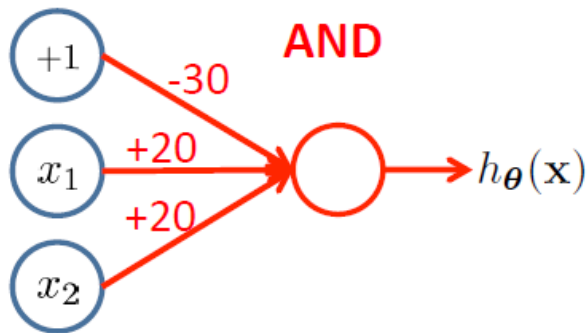
Boolean functions – OR

- What about OR
- $x_1, x_2 \in \{0,1\}$
- $y = x_1 \text{ OR } x_2$

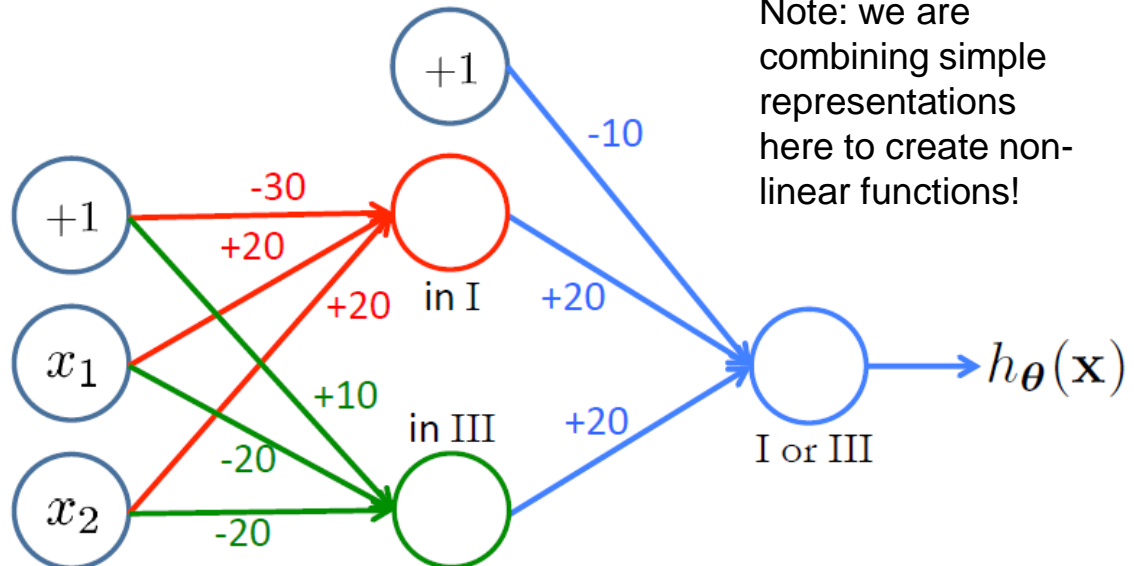



Boolean functions – NOT XOR

- What about NOT XOR: $y = NOT (x_1 XOR x_2)$



NOT XOR




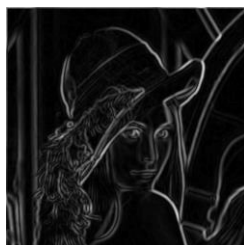
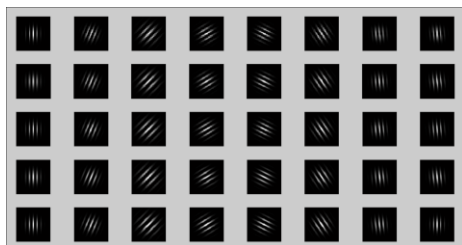
Note: we are combining simple representations here to create non-linear functions!

Feature learning: images



Traditionally:


Write a set of feature extractors:
e.g. texture and edge detectors




Insert into
a classifier

Neural network:

Feed raw pixels
into network



Each layer has
increasingly
abstract
representations



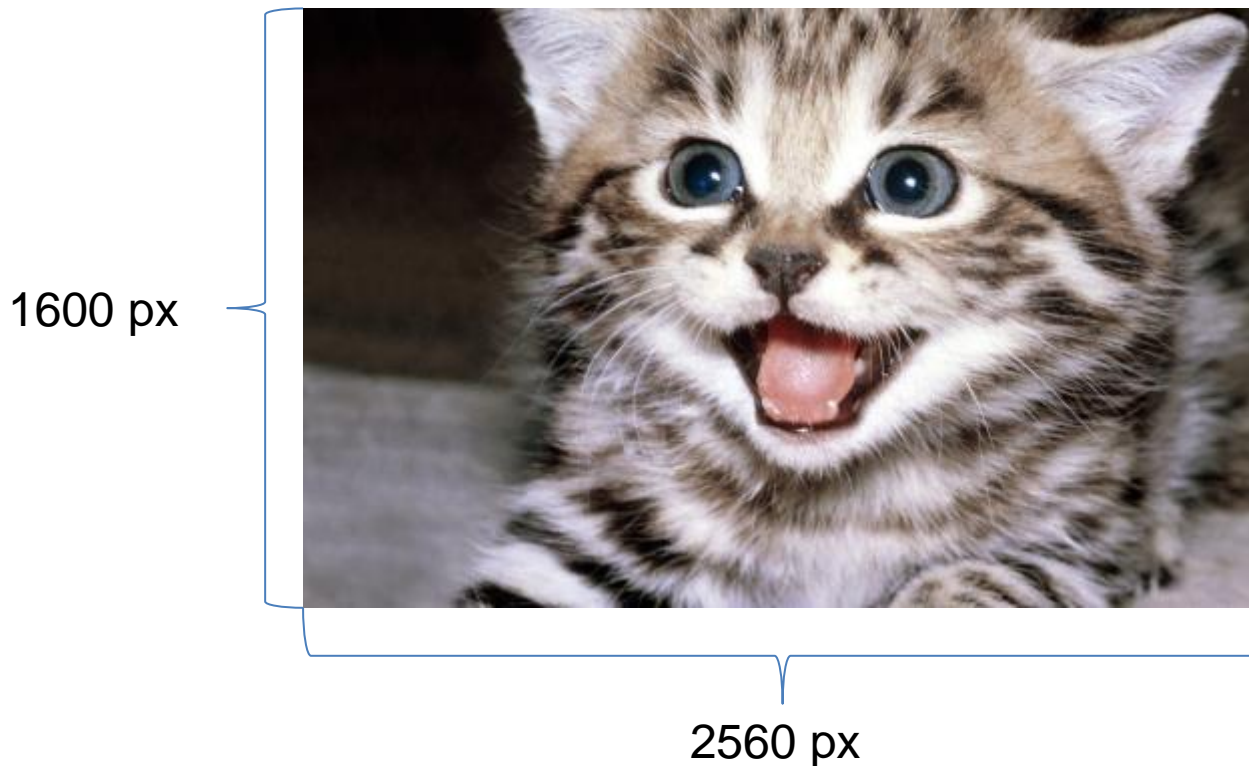
Final layer classifies

Convolutional neural networks

Working with images → too many weights!

(input neuron pre pixel)

ConvNets/CNNs: learn small patterns of shared weights



Convolutional Nets

- Use the same set of weights (a small patch – the filter) throughout the image
- Learn the weights \rightarrow feature detectors



Convolutional Nets

- Use the same set of weights (a small patch – the filter) throughout the image
- Learn the weights \rightarrow feature detectors



Convolutional Nets

- Use the same set of weights (a small patch – the filter) throughout the image
- Learn the weights \rightarrow feature detectors



Convolutional Nets

- Use the same set of weights (a small patch – the filter) throughout the image
- Learn the weights \rightarrow feature detectors




Convolutional Nets

- Use the same set of weights (a small patch – the filter) throughout the image
- Learn the weights \rightarrow feature detectors




Convolutional Nets

- Learn multiple feature detectors
- Corresponding maps indicate presence of features




Convolutional Nets

- Learn multiple feature detectors
- Corresponding maps indicate presence of features




Convolutional Nets

- Learn multiple feature detectors
- Corresponding maps indicate presence of features



Convolutional Nets

- Learn multiple feature detectors
- Corresponding maps indicate presence of features




Convolutional Nets

- Learn multiple feature detectors
- Corresponding maps indicate presence of features




Convolutional Nets



Convolutional Nets

- Repeat process on next layer




Convolutional Nets

Layer 2

Layer 1

Input




Convolutional Nets

Layer 3


Layer 2

Layer 1

Input



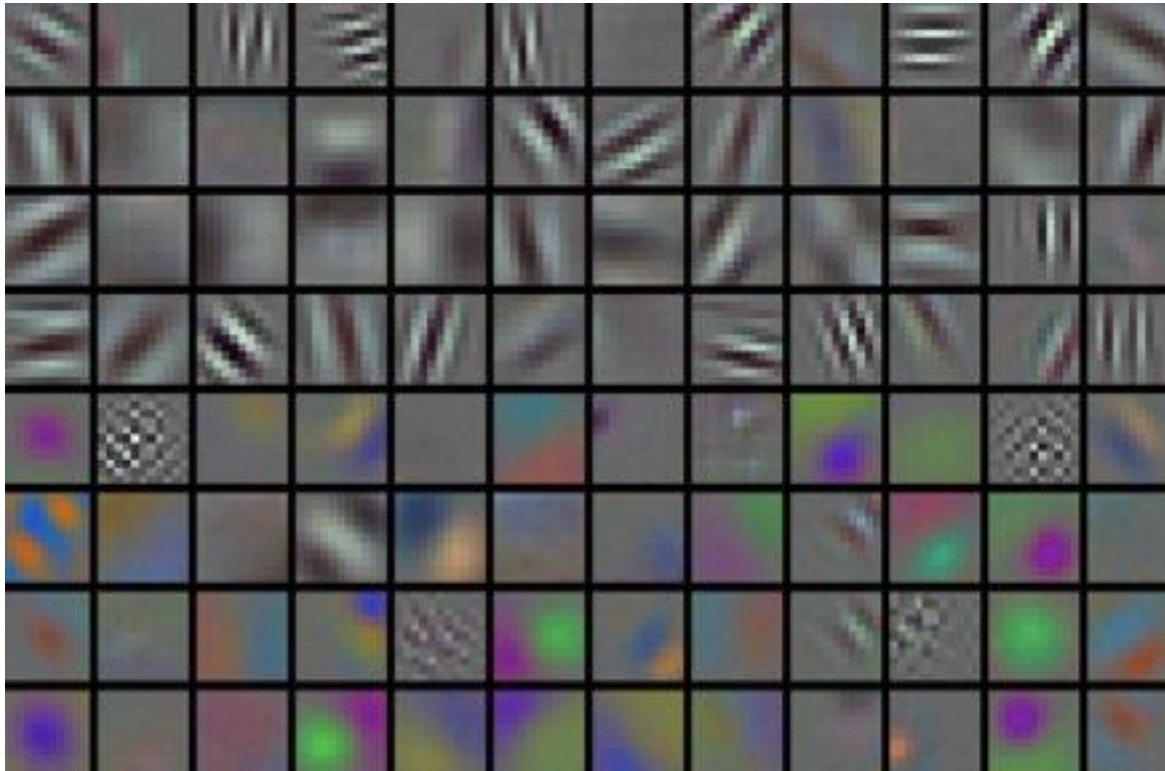
Convolutional Nets



Think of all
the data
needed to
train this!

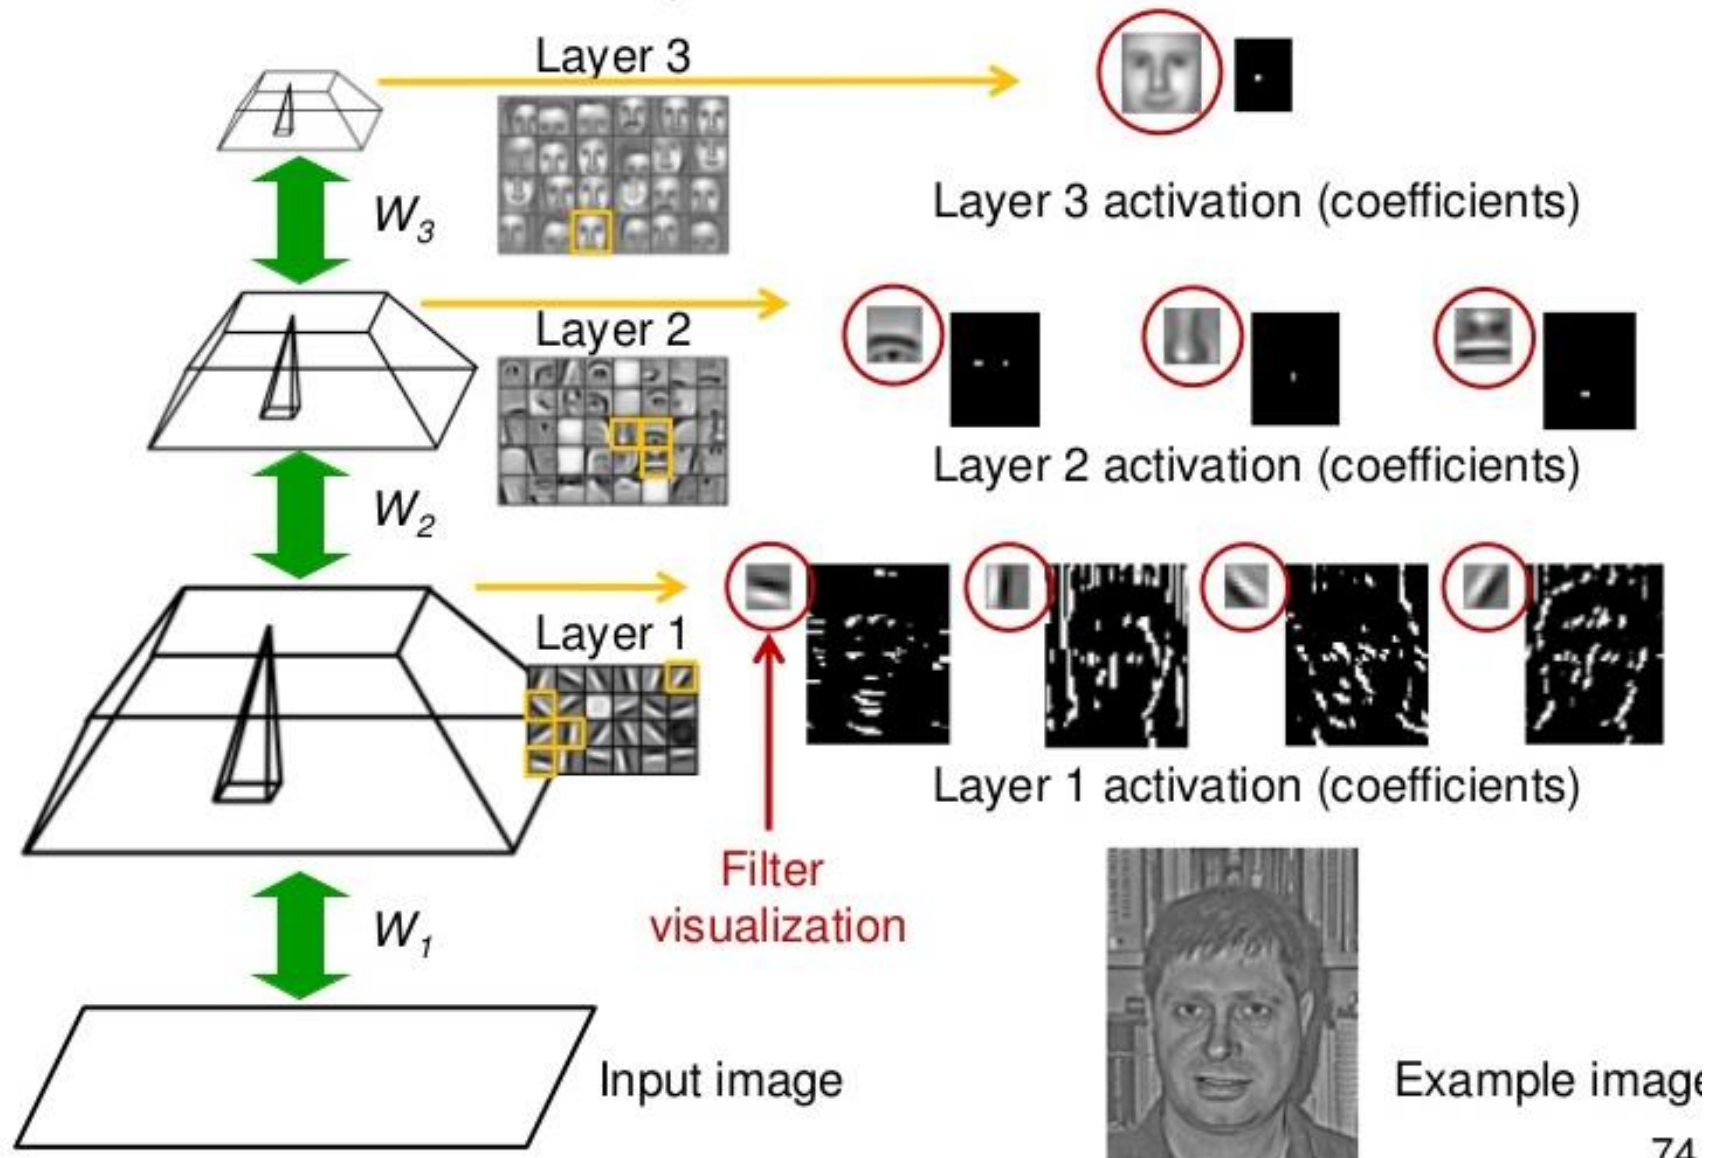
Extracted features

- Features (filters) in the first hidden layer



Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. "Imagenet classification with deep convolutional neural networks." *Advances in neural information processing systems*. 2012.

Deep Learning: Stacked Features



Recap

- Brains as motivation
- Different types of neurons
 - Relation to linear regression, logistic regression, perceptrons
- Feed-forward networks
- Hidden neurons
- Multiclass classification
- Boolean functions
- Convolutional networks
- Next time: training a neural network